

## CS 338 Fun Assignment #5

### Question 1.

Based on the given database scenario and concurrent transactions, determine the appropriate isolation levels. Consider an online banking system with the following database schema:

- Account(account\_id, customer\_id, balance, account\_type)
- Transaction\_Log(log\_id, account\_id, transaction\_type, amount, timestamp)

For each of the following transaction scenarios, determine the **lowest appropriate isolation level** and explain your reasoning:

(a) Transaction T1 performs a single account balance update:

```
-- T1:
UPDATE Account
SET balance = balance + 500
WHERE account_id = 12345;
COMMIT;
```

(b) Transaction T2 generates a monthly account statement:

```
-- T2:
SELECT customer_id, balance, account_type
FROM Account
WHERE customer_id = 67890;

SELECT log_id, transaction_type, amount, timestamp
FROM Transaction_Log
WHERE account_id IN (
    SELECT account_id FROM Account WHERE customer_id = 67890
);
COMMIT;
```

(c) Transaction T3 calculates and updates interest for all savings accounts:

```
-- T3:
SELECT account_id, balance
FROM Account
WHERE account_type = 'SAVINGS';

-- For each account found, calculate 2% interest
UPDATE Account
SET balance = balance * 1.02
WHERE account_type = 'SAVINGS';

-- Log the interest transactions
INSERT INTO Transaction_Log
```

```
SELECT NEXTVAL('log_seq'), account_id, 'INTEREST', balance * 0.02, NOW()  
FROM Account  
WHERE account_type = 'SAVINGS';  
COMMIT;
```

## Question 2.

Conceptual questions on concurrency anomalies and storage search method.

(a) Explain the difference between an **unrepeatable read** and a **phantom read**.

(b) Suggest 1 **benefits** and 1 **potential problems** of maintaining secondary indexes in a high-update OnLine Transaction Processing workload (systems that handle a large number of short, fast transactions in real time, online banking for example).

### Question 3.

Design and analyze a database system that handles both transaction isolation and indexing requirements. You are designing a library management system with the following tables:

- Book(isbn, title, author, category, copies\_available)
- Member(member\_id, name, email, membership\_type)
- Checkout(checkout\_id, member\_id, isbn, checkout\_date, due\_date, returned\_date)

(a) The system needs to handle the following concurrent transaction. Determine the minimum isolation level required.

A member checks out a book

```
-- T_CHECKOUT:
SELECT copies_available FROM Book WHERE isbn = '978-0123456789';
-- If copies_available > 0:
UPDATE Book SET copies_available = copies_available - 1
WHERE isbn = '978-0123456789';
INSERT INTO Checkout
VALUES (NEXTVAL('checkout_seq'), 12345, '978-0123456789',
      CURRENT_DATE, CURRENT_DATE + 14, NULL);
COMMIT;
```

Consider that multiple members might try to check out the last copy of the same book simultaneously.

(b) Consider the following two queries that are frequently executed on the library management system:

T1:

```
SELECT * FROM Book WHERE isbn = '978-0123456789';
```

T2:

```
SELECT b.title, b.author, c.checkout_date, c.due_date
FROM Book b, Checkout c
WHERE b.isbn = c.isbn
      AND c.returned_date IS NULL
      AND c.member_id = 12345;
```

For each query, design an optimal indexing strategy.